

# Fine-tuning vs. Model Size vs. Parameter-Efficient Tuning (LoRA): A Compact Experimental Report on Efficacy, Efficiency, and Reasoning Control

Alberto Rodero\*      Pablo Lobato\*

September 2025

## 1 Introduction

We evaluate the trade-offs between **full fine-tuning** (NoPeft), **LoRA**-based tuning with different ranks, and **base model size** using Qwen3-0.6B and Qwen3-1.7B on three task families: ARC (multiple-choice), OpenMath (numeric reasoning; lower is better), and SQuAD v2 (extractive QA). Additionally, we consider the model’s “**reasoning**” **mode** (a run-time behavior that may or may not be engaged depending on the generation configuration). In these runs, *reasoning was not locked*, which can inject variance into both latency and accuracy. This report answers practical questions about whether, when, and how to fine-tune; how LoRA rank matters; how results compare to a larger base model; and how to set reasoning going forward.

### Key questions.

- **Q1:** Is fine-tuning worth it in efficacy and efficiency?
- **Q2:** How does a fine-tuned 0.6B compare to a larger 1.7B base?
- **Q3:** If VRAM is abundant, is LoRA still worth using?
- **Q4:** Does LoRA rank materially affect outcomes?
- **Q5:** Partial vs. full tuning: LoRA vs. NoPeft?
- **Q6:** Cross-task transfer effects?
- **Q7:** If VRAM forces LoRA, is fine-tuning still worth it?

## 2 Experimental Setup

**Models.** Qwen3-0.6B (fine-tuned on: ARC, OpenMath, SQuAD; with NoPeft and LoRA ranks {32, 64, 256, 512, 1024}), Qwen3-0.6B base, Qwen3-1.7B base.

**Tasks & metrics.** ARC: macro-F1 (higher is better). OpenMathInstruct-2: average absolute difference (lower is better). SQuAD v2: F1 (higher is better). We also record mean latency (seconds) per task.

**Reasoning.** No explicit on/off control was enforced during these runs, so the model may have engaged/disengaged internal reasoning opportunistically. We treat this as a confound we will control in follow-up (Sec. ??).

---

\*Equal contribution

### 3 Full Results Table

Notes: “Math AbsDiff↓” is lower=better. Latencies are seconds. “Train Dataset” is the fine-tuning source (“\_base” for none).

Table 1: Full results across tasks. Lower is better for Math AbsDiff.

Notes: Latency in seconds. “Train DS” is the fine-tuning source (“\_base” = none).

	ARC F1	ARC Lat (s)	Math AbsDiff↓	Math Lat (s)	SQuAD F1	SQuAD Lat (s)	Train DS
Qwen3-0.6B-arc.SFT_None.Lora1024	0.4921	0.1557	22,871	0.4374	8.59	0.1939	arc
Qwen3-0.6B-arc.SFT_None.Lora512	0.4861	0.1570	22,871	0.4312	8.59	0.1955	arc
Qwen3-0.6B-arc.SFT_None.Lora256	0.4921	0.1549	22,871	0.4191	8.59	0.1948	arc
Qwen3-0.6B-arc.SFT_None.Lora64	0.4937	0.1601	22,871	0.1811	8.09	0.1952	arc
Qwen3-0.6B-arc.SFT_None.Lora32	0.4880	0.1595	22,871	0.4439	8.59	0.1958	arc
Qwen3-0.6B-arc.SFT_NoPeft_NoQuant	0.4905	0.0803	23,108	0.2137	18.89	0.2003	arc
Qwen3-0.6B-openmath.SFT_None.Lora1024	0.4990	0.9257	23,919	1.5384	8.48	0.2091	openmath
Qwen3-0.6B-openmath.SFT_None.Lora256	0.5031	0.8982	23,655	1.5016	8.48	0.1934	openmath
Qwen3-0.6B-openmath.SFT_None.Lora32	0.5095	0.8702	23,919	1.5776	8.48	0.1963	openmath
Qwen3-0.6B-openmath.SFT_NoPeft_NoQuant	0.5171	0.0605	16,540	0.0482	7.40	0.2439	openmath
Qwen3-0.6B-squad.SFT_None.Lora1024	0.5024	0.3178	23,647	2.1285	9.59	0.1951	squad
Qwen3-0.6B-squad.SFT_None.Lora256	0.5024	0.3116	23,647	2.0291	9.59	0.1908	squad
Qwen3-0.6B-squad.SFT_None.Lora32	0.5024	0.3095	23,523	1.9676	9.59	0.1950	squad
Qwen3-0.6B-squad.SFT_NoPeft_NoQuant	0.4542	0.1903	22,997	0.2203	27.95	0.2202	squad
Qwen3-0.6B_base	0.4932	1.1397	24,834	6.0139	10.07	0.2274	_base
Qwen3-1.7B_base	0.7986	3.4025	742	12.5197	30.36	0.2837	_base

### 4 Results by Question & Interpretation

#### Q1. Is fine-tuning worth it (efficacy & efficiency)?

Yes, when aligned to the target task, fine-tuning yields large quality gains and often lower latency.

- **OpenMath SFT (NoPeft):** Math abs diff improves by **33.4%** vs 0.6B\_base; ARC macro-F1 rises **+4.8%**; SQuAD F1 drops **-26.5%**. Latency massively drops on ARC (**-94.7%**) and Math (**-99.2%**), small increase on SQuAD (**+7.3%**).
- **SQuAD SFT (NoPeft):** SQuAD F1 jumps **+177.6%**; ARC macro-F1 dips **-7.9%**; Math improves mildly (**+7.4%** abs-diff reduction). Latency generally improves (ARC **-83.3%**, Math **-96.3%**, SQuAD **-3.2%**).

*Why?* Full-task alignment amplifies the relevant capabilities and stabilizes decoding behavior; our pipeline also appears to run tuned checkpoints more efficiently.

#### Q2. Fine-tuned 0.6B vs. larger 1.7B base?

**The 1.7B base dominates on quality but is slower.** ARC macro-F1 +61.9%, Math abs diff ~97% better, and SQuAD F1 +201.5% vs 0.6B\_base. Latency worsens markedly (ARC +198.5%, Math +108.2%, SQuAD +24.8%). **If latency/throughput matters, 0.6B + targeted SFT is the speed/price sweet spot.**

#### Q3. If VRAM is not a problem, should we still use LoRA?

**No. Prefer full fine-tuning.** In these runs, **NoPeft is both more accurate and often faster** than LoRA (Table 2). The only place where NoPeft is slightly slower is SQuAD (+0.029s), but it delivers a **+18.36** absolute F1 jump over the best LoRA.

Table 2: Best LoRA vs. NoPeft by task.

Task	Best LoRA model	Best LoRA score	Best LoRA lat (s)	NoPeft model	NoPeft score	NoPeft lat (s)
ARC	Qwen3-0.6B-arc_SFT_None_Lora64	0.4937	0.1601	Qwen3-0.6B-arc_SFT_NoPeft_NoQuant	0.4905	0.0803
OpenMath (abs diff ↓)	Qwen3-0.6B-openmath_SFT_None_Lora256	23.655	1.5016	Qwen3-0.6B-openmath_SFT_NoPeft_NoQuant	16.540	0.0482
SQuAD	Qwen3-0.6B-squad_SFT_None_Lora1024	9.59	0.1951	Qwen3-0.6B-squad_SFT_NoPeft_NoQuant	27.95	0.2202

#### Q4. Does LoRA rank matter?

**Little to no monotonic benefit from increasing rank.** Correlation between rank and score (LoRA-only subsets): **ARC F1**:  $\rho=0.021$ ; **OpenMath abs diff**:  $\rho=0.301$  (higher rank slightly worse); **SQuAD F1**: constant across ranks (no signal). Rank vs. latency has weak-to-moderate correlations (e.g., ARC  $\rho=-0.646$  suggests slightly lower latency at higher ranks, but absolute differences are tiny).

Table 3: Correlation between LoRA rank and performance/latency (LoRA-only).

Task	Corr(rank, score)	Corr(rank, latency)
ARC (F1)	0.021	-0.646
OpenMath (AbsDiff↓)	0.301	-0.233
SQuAD (F1)	N/A	0.321

*Why might rank not matter much?* With limited data or strongly structured tasks, the low-rank subspace often captures the critical adaptations; beyond a point, extra capacity (higher rank) faces diminishing returns and optimization noise. Also, our decoding and “reasoning” variability likely masks small rank effects.

#### Q5. LoRA (partial) vs. NoPeft (full) training?

**Full fine-tuning wins decisively.** OpenMath: NoPeft improves error by **7,115** absolute over best LoRA and is **~1.45 s** faster. SQuAD: NoPeft improves F1 by **+18.36** with only **+0.029 s** latency penalty. ARC: NoPeft is **~2× faster** than best LoRA, with essentially tied F1.

#### Q6. Cross-task transfer?

**Positive and negative transfer exist; pick your anchors carefully.**

- **OpenMath SFT (NoPeft)** → ARC: **+4.8%** (helpful), SQuAD: **−26.5%** (harmful).
- **SQuAD SFT (NoPeft)** → ARC: **−7.9%** (harmful), Math: **+7.4%** (helpful).
- **ARC SFT (NoPeft)** → SQuAD: **+87.6%** (surprisingly helpful), Math: **+7.0%**; ARC itself: **−0.5%** (negligible).

*Interpretation.* Extractive QA and numeric reasoning compete for capacity in different ways; tuning to one can suppress behaviors useful to the other. ARC SFT (full) seems to improve general reading/selection that transfers to SQuAD, while OpenMath SFT strengthens procedural reasoning that can conflict with extractive behavior.

#### Q7. If VRAM forces LoRA, is it still worth it?

**Yes, but temper expectations.** OpenMath LoRA-256 reduces error by **4.8%** vs 0.6B\_base (useful), but SQuAD LoRA variants lose **~4.8%** F1 vs 0.6B\_base, and ARC LoRA gains are marginal (**+0.1%**). **Recommendation:** When constrained, **LoRA  $\approx$  256** is a good default; otherwise prefer NoPeft.




Figure 1: Efficacy vs Latency

## 5 Limitations

**Reasoning mode was uncontrolled (confound).** **Single-seed/shot artifacts** may exist. **Latency** reflects this pipeline/hardware; other stacks may differ. No calibration/temperature sweeps are reported here.

## 6 Conclusions & Takeaways

- (1) Full fine-tuning beats LoRA on both accuracy and, in these runs, speed.
- (2) A larger base model is best on accuracy but slower; 0.6B+SFT is the speed/value sweet spot.
- (3) LoRA rank has weak, non-monotonic effects; use  $\sim 256$  by default when constrained.